


AI Resume Analyzer

Upload Resumes

Select PDF or plain text files to analyze.



Drag & Drop Files or Click to Browse
PDF, TXT (Max 5MB per file)

No files selected.

Process Resumes

Overview

This document outlines the design and development requirements for a Resume Analyzer Chatbot focused on job matching. The core system will ingest and index multiple resume files (PDF/plain text), accept a job description, and leverage AI to identify the most relevant candidates from the indexed pool. It will provide insights and answer queries in a conversational manner. Use of free-tier cloud services and open-source libraries is emphasized to build a robust and cost-effective solution. **Candidates may use an open-source resume dataset (e.g., public academic or community datasets).** A web-based user interface is optional and can be implemented as a bonus feature.

Core Functionality (Required)

1. Resume Ingestion and Processing

- **Index multiple resumes:** Ingest and process multiple resume files (PDF and plain text).
- **Parse content:** Extract text from resumes.
- **Chunk text:** Split parsed text into meaningful chunks (100–500 tokens with overlap).
- **Generate embeddings:** Produce vector embeddings for each chunk using a free-tier model.

- *Recommended:* Google Gemini (free tier, e.g., 1.5 Flash)
- *Other:* Hugging Face all-MiniLM-L6-v2

2. Vector Database Integration

- **Store embeddings:** Save generated embeddings in a vector database.
- **Efficient retrieval:** Implement top-K similarity search for queries.
 - *Recommended:* Aiven PostgreSQL with pgvector (free tier)
 - *Other:* Pinecone (free tier)

3. Retrieval-Augmented Generation (RAG) Chatbot

- **Job description input:** Accept a text description of the job role.
 - **Retrieve relevant chunks:** Perform vector search to find top resume chunks.
 - **LLM-based matching:** Use an LLM to generate conversational answers about candidate fit, citing retrieved content.
 - **Conversational interface:** Support follow-up questions for deeper insights.
 - *Recommended LLM:* Google Gemini (free tier)
 - *Other:* Open-source LLM (e.g., LLaMA) running locally
-

Optional Bonus Features

A. Web User Interface

- **Frontend (Bonus):** Implement a simple web UI for file upload, job description entry, and chat interaction.
 - *Tech Stack:* Next.js with TypeScript
 - *Hosting:* Vercel (Hobby/Free Tier)

B. SQL-Based Metadata Search

- **Extract metadata:** Tag resumes with structured metadata (skills, titles, experience).
 - **Metadata storage:** Save tags in a relational database (e.g., PostgreSQL alongside vector store).
 - **Metadata API:** Expose an endpoint for SQL queries (e.g., `SELECT * FROM resume_metadata WHERE skills @> ARRAY['TypeScript'] AND years_experience >= 5`).
-

Technical Stack & Free-Tier Considerations

Candidates should choose and justify one option per component, demonstrating mitigation of free-tier limits.

- **Backend:** Python (FastAPI/Flask) or TypeScript (Next.js API routes)
 - *Hosting:* Render (Python) or Vercel (Next.js)
 - **Object Storage:** Cloudflare R2 or Vercel Blob
 - **Orchestration & Processing:** LangChain for loading, splitting, embedding, RAG orchestration, and metadata extraction
-

Deliverables

Required

1. **Backend System:**
 - Batch resume parsing & chunking
 - Embedding generation & storage
 - Vector similarity search for job queries
 - RAG-based response generation
2. **Public GitHub Repository:**
 - README with setup, usage, free-tier notes
 - Clean, modular, documented code

Optional (Bonus)

- **Deployed Web Application:**
 - File upload UI, job description input, chat interface
 - Live URL (e.g., Vercel)
 - **Metadata Search API:**
 - Endpoint for structured SQL queries
 - **Demo Walkthrough:**
 - Short video or write-up showcasing features
-

Notes for Candidates: Focus first on delivering the core ingestion, retrieval, and RAG functionality. If time permits, implement the optional web UI and metadata search for extra credit.